

Basic Data Analysis Tutorial

August 31, 2023

1 Basic Data Analysis Tutorial

```
[4]: # Install the libraries
# !pip install skillsnetwork
# !pip install matplotlib
# !pip install numpy
# !pip install pandas
# !pip install seaborn
```

```
[57]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import skillsnetwork
import warnings
```

When engaged in data analysis, having a clear understanding of your objective is essential.

Let's analyze the 'price'!

Import the dataset

```
[4]: path = "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
↳IBMDeveloperSkillsNetwork-DA0101EN-SkillsNetwork/labs/Data%20files/auto.csv"
df = pd.read_csv(path, header = None)
```

2 DATA EXPLORATION

```
[5]: df.head(5)
```

```
[5]:
```

	0	1	2	3	4	5	6	7	8	9	...	\
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	

	16	17	18	19	20	21	22	23	24	25
0	130	mpfi	3.47	2.68	9.0	111	5000	21	27	13495
1	130	mpfi	3.47	2.68	9.0	111	5000	21	27	16500
2	152	mpfi	2.68	3.47	9.0	154	5000	19	26	16500
3	109	mpfi	3.19	3.40	10.0	102	5500	24	30	13950
4	136	mpfi	3.19	3.40	8.0	115	5500	18	22	17450

[5 rows x 26 columns]

Visualize the first 5 rows of the dataframe

```
[6]: df.describe()
```

```
[6]:
```

	0	9	10	11	12 \
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	0.834146	98.756585	174.049268	65.907805	53.724878
std	1.245307	6.021776	12.337289	2.145204	2.443522
min	-2.000000	86.600000	141.100000	60.300000	47.800000
25%	0.000000	94.500000	166.300000	64.100000	52.000000
50%	1.000000	97.000000	173.200000	65.500000	54.100000
75%	2.000000	102.400000	183.100000	66.900000	55.500000
max	3.000000	120.900000	208.100000	72.300000	59.800000

	13	16	20	23	24
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	2555.565854	126.907317	10.142537	25.219512	30.751220
std	520.680204	41.642693	3.972040	6.542142	6.886443
min	1488.000000	61.000000	7.000000	13.000000	16.000000
25%	2145.000000	97.000000	8.600000	19.000000	25.000000
50%	2414.000000	120.000000	9.000000	24.000000	30.000000
75%	2935.000000	141.000000	9.400000	30.000000	34.000000
max	4066.000000	326.000000	23.000000	49.000000	54.000000

The describe() function provides the summary statistics of the dataset.

(Only contains numerical attributes)

```
[7]: df.describe(include = 'all')
```

```
[7]:
```

	0	1	2	3	4	5	6	7	8 \
count	205.000000	205	205	205	205	205	205	205	205
unique	NaN	52	22	2	2	3	5	3	2
top	NaN	?	toyota	gas	std	four	sedan	fwd	front
freq	NaN	41	32	185	168	114	96	120	202
mean	0.834146	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
std	1.245307	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
min	-2.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
25%	0.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

50%	1.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
75%	2.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
max	3.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

	9	...	16	17	18	19	20	21	22	\
count	205.000000	...	205.000000	205	205	205	205.000000	205	205	
unique	NaN	...	NaN	8	39	37	NaN	60	24	
top	NaN	...	NaN	mpfi	3.62	3.40	NaN	68	5500	
freq	NaN	...	NaN	94	23	20	NaN	19	37	
mean	98.756585	...	126.907317	NaN	NaN	NaN	10.142537	NaN	NaN	
std	6.021776	...	41.642693	NaN	NaN	NaN	3.972040	NaN	NaN	
min	86.600000	...	61.000000	NaN	NaN	NaN	7.000000	NaN	NaN	
25%	94.500000	...	97.000000	NaN	NaN	NaN	8.600000	NaN	NaN	
50%	97.000000	...	120.000000	NaN	NaN	NaN	9.000000	NaN	NaN	
75%	102.400000	...	141.000000	NaN	NaN	NaN	9.400000	NaN	NaN	
max	120.900000	...	326.000000	NaN	NaN	NaN	23.000000	NaN	NaN	

	23	24	25
count	205.000000	205.000000	205
unique	NaN	NaN	187
top	NaN	NaN	?
freq	NaN	NaN	4
mean	25.219512	30.751220	NaN
std	6.542142	6.886443	NaN
min	13.000000	16.000000	NaN
25%	19.000000	25.000000	NaN
50%	24.000000	30.000000	NaN
75%	30.000000	34.000000	NaN
max	49.000000	54.000000	NaN

[11 rows x 26 columns]

By using the command `describe(include='all')`, it provides summary statistics for all attributes, including: - `unique`: indicating how many different unique 'names' or 'elements' are within each categorical attribute. - `top`: representing the element that appeared most frequently among the categorical attributes. - `freq`: denoting the frequency or the number of times it appeared.

[8]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column  Non-Null Count  Dtype
---  -
0    0      205 non-null    int64
1    1      205 non-null    object
2    2      205 non-null    object
3    3      205 non-null    object
```

```

4  4      205 non-null    object
5  5      205 non-null    object
6  6      205 non-null    object
7  7      205 non-null    object
8  8      205 non-null    object
9  9      205 non-null    float64
10 10     205 non-null    float64
11 11     205 non-null    float64
12 12     205 non-null    float64
13 13     205 non-null    int64
14 14     205 non-null    object
15 15     205 non-null    object
16 16     205 non-null    int64
17 17     205 non-null    object
18 18     205 non-null    object
19 19     205 non-null    object
20 20     205 non-null    float64
21 21     205 non-null    object
22 22     205 non-null    object
23 23     205 non-null    int64
24 24     205 non-null    int64
25 25     205 non-null    object
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB

```

The `info()` function provides us with the data type of each attribute, as well as helps in identifying null values.

```
[9]: df.dtypes
```

```

[9]: 0      int64
1      object
2      object
3      object
4      object
5      object
6      object
7      object
8      object
9      float64
10     float64
11     float64
12     float64
13     int64
14     object
15     object
16     int64
17     object

```

```

18     object
19     object
20    float64
21     object
22     object
23      int64
24      int64
25     object
dtype: object

```

Alternatively, we can use dtype.

3 We noticed that there are several things we need to fix

- 1.) We noticed that this dataframe does not have a header
- 2.) We noticed there are missing values in the dataset
- 3.) We noticed that the some data types do not match

Let's fix them!

```
[10]: df.replace('?', np.nan, inplace = True)
```

Replacing all '?' to NaN to make it easier

4 1.) Adding Header

```
[11]: headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration",
↳ "num-of-doors", "body-style",
      "drive-wheels", "engine-location", "wheel-base",
↳ "length", "width", "height", "curb-weight", "engine-type",
      "num-of-cylinders",
↳ "engine-size", "fuel-system", "bore", "stroke", "compression-ratio", "horsepower",
      "peak-rpm", "city-mpg", "highway-mpg", "price"]
df.columns = headers
df.head()
```

```
[11]:  symboling  normalized-losses      make fuel-type aspiration num-of-doors \
0         3             NaN  alfa-romero    gas      std         two
1         3             NaN  alfa-romero    gas      std         two
2         1             NaN  alfa-romero    gas      std         two
3         2             164      audi     gas      std         four
4         2             164      audi     gas      std         four

      body-style drive-wheels engine-location  wheel-base  ...  engine-size \
0  convertible         rwd         front         88.6  ...         130
1  convertible         rwd         front         88.6  ...         130

```

2	hatchback	rwd	front	94.5	...	152
3	sedan	fwd	front	99.8	...	109
4	sedan	4wd	front	99.4	...	136

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

5 2.) Dealing With Missing Values

```
[12]: missing_data = df.isnull()

# This forloop check the number of missing values for each attribute:

# for column in missing_data.columns.values.tolist():
#     print(column)
#     print(missing_data[column].value_counts())
#     print('')
```

- “normalized-losses”: 41 missing data
- “num-of-doors”: 2 missing data
- “bore”: 4 missing data
- “stroke” : 4 missing data
- “horsepower”: 2 missing data
- “peak-rpm”: 2 missing data
- “price”: 4 missing data

Let’s deal with the missing values

There are three most common ways to deal with missing values: * 1.) Drop the missing values * 2.) Replace the missing values * 3.) Leave the missing values

```
[13]: # Drop the missing values in the subset price
df.dropna(subset = ['price'], inplace = True)
```

```
[14]: # Calculate the averages
avg_normalized_losses = df['normalized-losses'].astype('float').mean()
avg_bore = df['bore'].astype('float').mean()
avg_stroke = df['stroke'].astype('float').mean()
avg_peak_rpm = df['peak-rpm'].astype('float').mean()
avg_horsepower = df['horsepower'].astype('float').mean()

# Replace the missing value with the averages calculated above
df['normalized-losses'].replace(np.nan, avg_normalized_losses, inplace=True)
df['bore'].replace(np.nan, avg_bore, inplace=True)
df['stroke'].replace(np.nan, avg_stroke, inplace=True)
df['peak-rpm'].replace(np.nan, avg_peak_rpm, inplace=True)
df['horsepower'].replace(np.nan, avg_horsepower, inplace=True)
```

```
[15]: # For categorical attributes, we can see how many units of each variable we have
print(df['num-of-doors'].value_counts())
```

```
four    113
two      86
Name: num-of-doors, dtype: int64
```

```
[16]: # 'four' is the most common, so let's replace the missing values with it
# However, value_counts() only works on pandas series
# As a result, we only include one bracket df['drive-wheels'], not two brackets
↳ df[['drive-wheels']].
df['num-of-doors'].replace(np.nan, 'four', inplace=True)
```

6 3.) Fixing the Data Type

- 'normalized-losses' and 'horsepower' should be of type int
- 'bore', 'stroke', 'peak-rpm' and 'price' should be of type float

```
[17]: # Change to type int
df[['normalized-losses', 'horsepower']] = df[['normalized-losses',
↳ 'horsepower']].astype('int')

# Change to type float
df[['bore', 'stroke', 'peak-rpm', 'price']] = df[['bore', 'stroke', 'peak-rpm',
↳ 'price']].astype('float')
```

```
[18]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 201 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
#   ...
```

```

---  -----  -----  -----
0   symboling      201 non-null   int64
1   normalized-losses 201 non-null   int64
2   make           201 non-null   object
3   fuel-type      201 non-null   object
4   aspiration      201 non-null   object
5   num-of-doors    201 non-null   object
6   body-style      201 non-null   object
7   drive-wheels    201 non-null   object
8   engine-location 201 non-null   object
9   wheel-base      201 non-null   float64
10  length          201 non-null   float64
11  width           201 non-null   float64
12  height          201 non-null   float64
13  curb-weight     201 non-null   int64
14  engine-type      201 non-null   object
15  num-of-cylinders 201 non-null   object
16  engine-size      201 non-null   int64
17  fuel-system      201 non-null   object
18  bore            201 non-null   float64
19  stroke          201 non-null   float64
20  compression-ratio 201 non-null   float64
21  horsepower      201 non-null   int64
22  peak-rpm        201 non-null   float64
23  city-mpg        201 non-null   int64
24  highway-mpg     201 non-null   int64
25  price           201 non-null   float64
dtypes: float64(9), int64(7), object(10)
memory usage: 42.4+ KB

```

7 Data Transformation

Sometimes dataframe has a different format than what we need

Hence we need to transform data into a common format

```

[19]: # df.head()

[20]: # Let's transform mpg into L/100km!
      # L/100km = 235/mpg
      df['city-mpg'] = 235/df['city-mpg']
      df['highway-mpg'] = 235/df['highway-mpg']

      # Don't forget to rename the column name after the transformation!
      df.rename(columns={'city-mpg': 'city-L/100km', 'highway-mpg' : 'highway-L/
      ↪100km'}, inplace = True)

```



```
df[['city-L/100km', 'highway-L/100km']] ## select and display 'city-L/100km' and
↳ 'highway-L/100km'
```

```
[20]:
```

	city-L/100km	highway-L/100km
0	11.190476	8.703704
1	11.190476	8.703704
2	12.368421	9.038462
3	9.791667	7.833333
4	13.055556	10.681818
..
200	10.217391	8.392857
201	12.368421	9.400000
202	13.055556	10.217391
203	9.038462	8.703704
204	12.368421	9.400000

```
[201 rows x 2 columns]
```

8 Data Normalization

It is important to normalize our dataframe

Think about this scenario, we have a dataframe that contain an age attribute and an income attribute

```
[21]: data1 = {'Age' : [21,38,25,41],
              'Income' : [5000,5500,3000,7700]}
ex1 = pd.DataFrame(data1)
ex1
```

```
[21]:
```

	Age	Income
0	21	5000
1	38	5500
2	25	3000
3	41	7700

Income has a bigger effect on the analysis than age because it is larger, we need to make them have equal effects by normalizing them.

There are different ways you can perform normalization such as z-score

However, let's keep it easy and perform value/max_value

```
[22]: # Let's go back to our dataframe - df and perform this operation on the
↳ 'length', 'width' and 'height'

df['length'] = df['length']/df['length'].max()
df['width'] = df['width']/df['width'].max()
df['height'] = df['height']/df['height'].max()
```

```
df[['length', 'width', 'height']].head(6)
```

```
[22]:      length      width      height
0  0.811148  0.890278  0.816054
1  0.811148  0.890278  0.816054
2  0.822681  0.909722  0.876254
3  0.848630  0.919444  0.908027
4  0.848630  0.922222  0.908027
5  0.851994  0.920833  0.887960
```

9 Data Binning

Why binning?

We perform data binning to transform continuous numerical variables into discrete categorical 'bins' for better data analysis.

```
[23]: # First, we need the matplotlib library to help with data visualization
import matplotlib as plt
from matplotlib import pyplot

# Let's divide horsepower into 3 bins
# linspace(start_value, end_value, numbers_generated)
bins = np.linspace(min(df['horsepower']), max(df['horsepower']), 4)


# Assign group names
group_names = ['Low', 'Medium', 'High']

# Apply the cut function to determine what each value of df['horsepower']
↳ belongs to
df['horsepower-binned'] = pd.cut(df['horsepower'], bins, labels = group_names,
    ↳ include_lowest = True)
df[['horsepower', 'horsepower-binned']].head(10)
```

```
[23]:      horsepower horsepower-binned
0           111                Low
1           111                Low
2           154               Medium
3           102                Low
4           115                Low
5           110                Low
6           110                Low
7           110                Low
8           140               Medium
10          101                Low
```

```
[24]: # Visualize the 'bin'
pyplot.bar(group_names, df["horsepower-binned"].value_counts())
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[24]: Text(0.5, 1.0, 'horsepower bins')
```



Side note: by using the Altair library, visualization can be enhanced

10 Dummy Variables

We assign dummy variables to label categories. It is called ‘dummies’ because the numbers themselves don’t have inherent meanings.

```
[25]: # Let's assign dummy variables for 'fuel-type' which has two distinct values:
      ↪ 'gas' and 'diesel'
dummy_variable = pd.get_dummies(df['fuel-type'])

# Change the column name
```

```
dummy_variable.rename(columns = {'diesel' : 'fuel-type-diesel', 'gas' : 'fuel-type-gas'}, inplace = True)
```

```
[26]: dummy_variable.head()
```

```
[26]:    fuel-type-diesel  fuel-type-gas
0                0                1
1                0                1
2                0                1
3                0                1
4                0                1
```

```
[27]: # Merge the dataframe df and dummy_variable together
df = pd.concat([df,dummy_variable], axis = 1)

# Drop the original 'fuel-type' from the dataframe 'df'
df.drop('fuel-type', axis = 1, inplace = True)
```

11 Data Correlation

Correlation is extremely important when it comes to data analysis, it tells us whether the variables are interdependent.

If a variable have a strong correlation with another, it is an indicator that the variable is a good predictor.

```
[28]: # To see the correlations, simply use the function corr()
df.corr()
```

/tmp/ipykernel_450/2902987488.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
df.corr()
```

```
[28]:
```

	symboling	normalized-losses	wheel-base	length	\
symboling	1.000000	0.466264	-0.535987	-0.365404	
normalized-losses	0.466264	1.000000	-0.056661	0.019424	
wheel-base	-0.535987	-0.056661	1.000000	0.876024	
length	-0.365404	0.019424	0.876024	1.000000	
width	-0.242423	0.086802	0.814507	0.857170	
height	-0.550160	-0.373737	0.590742	0.492063	
curb-weight	-0.233118	0.099404	0.782097	0.880665	
engine-size	-0.110581	0.112360	0.572027	0.685025	
bore	-0.139896	-0.029800	0.493203	0.608941	
stroke	-0.007992	0.055127	0.157964	0.123913	
compression-ratio	-0.182196	-0.114713	0.250313	0.159733	

horsepower	0.075776	0.217300	0.371297	0.579688
peak-rpm	0.279719	0.239544	-0.360233	-0.286035
city-L/100km	0.066171	0.238567	0.476153	0.657373
highway-L/100km	-0.029807	0.181189	0.577576	0.707108
price	-0.082391	0.133999	0.584642	0.690628
fuel-type-diesel	-0.196735	-0.101546	0.307237	0.211187
fuel-type-gas	0.196735	0.101546	-0.307237	-0.211187

	width	height	curb-weight	engine-size	bore \
symboling	-0.242423	-0.550160	-0.233118	-0.110581	-0.139896
normalized-losses	0.086802	-0.373737	0.099404	0.112360	-0.029800
wheel-base	0.814507	0.590742	0.782097	0.572027	0.493203
length	0.857170	0.492063	0.880665	0.685025	0.608941
width	1.000000	0.306002	0.866201	0.729436	0.544879
height	0.306002	1.000000	0.307581	0.074694	0.180327
curb-weight	0.866201	0.307581	1.000000	0.849072	0.644041
engine-size	0.729436	0.074694	0.849072	1.000000	0.572516
bore	0.544879	0.180327	0.644041	0.572516	1.000000
stroke	0.188814	-0.060822	0.167412	0.205806	-0.055390
compression-ratio	0.189867	0.259737	0.156433	0.028889	0.001250
horsepower	0.614972	-0.086901	0.758001	0.822636	0.566786
peak-rpm	-0.245852	-0.309913	-0.279350	-0.256753	-0.267338
city-L/100km	0.673363	0.003811	0.785353	0.745059	0.554726
highway-L/100km	0.736728	0.084301	0.836921	0.783465	0.559197
price	0.751265	0.135486	0.834415	0.872335	0.543154
fuel-type-diesel	0.244356	0.281578	0.221046	0.070779	0.054435
fuel-type-gas	-0.244356	-0.281578	-0.221046	-0.070779	-0.054435

	stroke	compression-ratio	horsepower	peak-rpm \
symboling	-0.007992	-0.182196	0.075776	0.279719
normalized-losses	0.055127	-0.114713	0.217300	0.239544
wheel-base	0.157964	0.250313	0.371297	-0.360233
length	0.123913	0.159733	0.579688	-0.286035
width	0.188814	0.189867	0.614972	-0.245852
height	-0.060822	0.259737	-0.086901	-0.309913
curb-weight	0.167412	0.156433	0.758001	-0.279350
engine-size	0.205806	0.028889	0.822636	-0.256753
bore	-0.055390	0.001250	0.566786	-0.267338
stroke	1.000000	0.187854	0.097598	-0.063720
compression-ratio	0.187854	1.000000	-0.214392	-0.435721
horsepower	0.097598	-0.214392	1.000000	0.107882
peak-rpm	-0.063720	-0.435721	0.107882	1.000000
city-L/100km	0.036285	-0.299372	0.889454	0.115813
highway-L/100km	0.047199	-0.223361	0.840695	0.017736
price	0.082267	0.071107	0.809729	-0.101542
fuel-type-diesel	0.241033	0.985231	-0.168941	-0.475759
fuel-type-gas	-0.241033	-0.985231	0.168941	0.475759

	city-L/100km	highway-L/100km	price	fuel-type-diesel \
symboling	0.066171	-0.029807	-0.082391	-0.196735
normalized-losses	0.238567	0.181189	0.133999	-0.101546
wheel-base	0.476153	0.577576	0.584642	0.307237
length	0.657373	0.707108	0.690628	0.211187
width	0.673363	0.736728	0.751265	0.244356
height	0.003811	0.084301	0.135486	0.281578
curb-weight	0.785353	0.836921	0.834415	0.221046
engine-size	0.745059	0.783465	0.872335	0.070779
bore	0.554726	0.559197	0.543154	0.054435
stroke	0.036285	0.047199	0.082267	0.241033
compression-ratio	-0.299372	-0.223361	0.071107	0.985231
horsepower	0.889454	0.840695	0.809729	-0.168941
peak-rpm	0.115813	0.017736	-0.101542	-0.475759
city-L/100km	1.000000	0.958306	0.789898	-0.241282
highway-L/100km	0.958306	1.000000	0.801118	-0.158091
price	0.789898	0.801118	1.000000	0.110326
fuel-type-diesel	-0.241282	-0.158091	0.110326	1.000000
fuel-type-gas	0.241282	0.158091	-0.110326	-1.000000

	fuel-type-gas
symboling	0.196735
normalized-losses	0.101546
wheel-base	-0.307237
length	-0.211187
width	-0.244356
height	-0.281578
curb-weight	-0.221046
engine-size	-0.070779
bore	-0.054435
stroke	-0.241033
compression-ratio	-0.985231
horsepower	0.168941
peak-rpm	0.475759
city-L/100km	0.241282
highway-L/100km	0.158091
price	-0.110326
fuel-type-diesel	-1.000000
fuel-type-gas	1.000000

```
[29]: # Let's the explore the correlation between 'bore', 'stroke', 'peak-rpm' and
      ↪ 'horsepower'
      df[['bore', 'stroke', 'peak-rpm', 'horsepower']].corr()
```


```
[29]:      bore      stroke  peak-rpm  horsepower
bore      1.000000 -0.055390 -0.267338      0.566786
```

stroke	-0.055390	1.000000	-0.063720	0.097598
peak-rpm	-0.267338	-0.063720	1.000000	0.107882
horsepower	0.566786	0.097598	0.107882	1.000000

```
[30]: # To see the correlation between 'horsepower' and 'price' visually
```

```
sns.regplot(x= 'horsepower', y = 'price', data= df)
```

```
[30]: <Axes: xlabel='horsepower', ylabel='price'>
```



```
[31]: df[['horsepower', 'price']].corr()
```

```
[31]:
```

	horsepower	price
horsepower	1.000000	0.809729
price	0.809729	1.000000


Above, we can see there is a positive correlation relationship between the variable 'horsepower' and 'price'.

Also, with a correlation of 0.809729. It seems like a pretty good predictor.

Did you notice we are only dealing with continuous numerical variables? What about categorical variables?

```
[32]: #To determine whether a categorical variable is a good indicator visually, we
      ↪ can use boxplot.
      sns.boxplot(x= 'engine-location', y= 'price', data= df)
```

```
[32]: <Axes: xlabel='engine-location', ylabel='price'>
```



The distinct differences between the engine-location shows it might be a good indicator of price.

If they share significant overlap, it might not be a good indicator. However, this might be too general and deeper analysis is usually required.

```
[33]: # Let's revisit the value_counts() function again for categorical variables.
      # It counts the number of units and covert to a dataframe and change the column
      ↪ name
      drive_wheels_counts = df['drive-wheels'].value_counts().to_frame()
      drive_wheels_counts.rename(columns={'drive-wheels': 'value_counts'},
      ↪ inplace=True)
      drive_wheels_counts
```

```
[33]:      value_counts
      fwd          118
      rwd           75
```


4wd

8

```
[34]: # Renme the index to 'drive-wheels'
drive_wheels_counts.index.name = 'drive-wheels'
drive_wheels_counts
```

```
[34]:          value_counts
drive-wheels
fwd          118
rwd           75
4wd           8
```

12 Data Grouping

We use the ‘groupby’ method to group the data by different categories.

Let’s do an example and group by the variable ‘drive-wheels’!

```
[35]: # Select the coumms 'drive-wheels', 'body-style' and 'price'.
df_group_one = df[['drive-wheels', 'body-style', 'price']]

# Group and calculate the average price for each of the different categories
df_group_one = df_group_one.groupby(['drive-wheels', 'body-style'], as_index =_
↪False).mean()
df_group_one
```

```
[35]:   drive-wheels  body-style      price
0          4wd    hatchback  7603.000000
1          4wd      sedan    12647.333333
2          4wd      wagon    9095.750000
3          fwd  convertible  11595.000000
4          fwd    hardtop    8249.000000
5          fwd    hatchback   8396.387755
6          fwd      sedan    9811.800000
7          fwd      wagon    9997.333333
8          rwd  convertible  23949.600000
9          rwd    hardtop   24202.714286
10         rwd    hatchback  14337.777778
11         rwd      sedan   21711.833333
12         rwd      wagon   16994.222222
```

Pivot

We can make it into a pivot table so it is easier to visualize.

```
[36]: grouped_pivot = df_group_one.pivot(index= 'drive-wheels', columns= 'body-style')
grouped_pivot
```

```
# Often we will encounter missing values, so let's fill it with 0
grouped_pivot.replace(np.nan, 0, inplace = True)
# OR
# grouped_pivot = grouped_pivot.fillna(0)
grouped_pivot
```

```
[36]:
```

	price			
body-style	convertible	hardtop	hatchback	sedan
drive-wheels				
4wd	0.0	0.000000	7603.000000	12647.333333
fwd	11595.0	8249.000000	8396.387755	9811.800000
rwd	23949.6	24202.714286	14337.777778	21711.833333

body-style	wagon
drive-wheels	
4wd	9095.750000
fwd	9997.333333
rwd	16994.222222

13 Model - Linear Regression Model

The linear regression model help us understand the relationship between two variables, the dependent(Y) and the independent(X) variable.

Linear Function

$$\hat{Y} = a + bX$$

- where 'a' is the intercept and 'b' is the slope

```
[38]: from sklearn.linear_model import LinearRegression
```

Let's perform the linear regression model for engine size and price.

```
[39]: lm = LinearRegression()
X = df[['engine-size']] #dataframe
Y = df['price'] #series

#fit the linear model
lm.fit(X,Y)
```

```
[39]: LinearRegression()
```

```
[40]: # Calculate the intercept
lm.intercept_
```

```
[40]: -7963.338906281046
```

```
[41]: # Calculate the slope
lm.coef_
```

```
[41]: array([166.86001569])
```

```
[42]: # <b>Therefore, we have:</b>

# $$
#  $\hat{Y} = -7963.34 + 166.86 * X$ 
# $$
```

14 Model - Multiple Linear Regression Model

What if we want to examine the relationship between the dependent variable with more than one predictor variables?

$$\hat{Y} = a + b_1X_1 + b_2X_2 + b_3X_3 + b_4X_4$$

```
[44]: # Let's explore the relationships of 'Horsepower', 'Curb-weight', 'Engine-size'
# and 'Highway-L/100km' with 'price'
lm1 = LinearRegression()
Z = df[['horsepower', 'curb-weight', 'engine-size', 'highway-L/100km']]

# Fit the multiple linear model
lm1.fit(Z,Y)
```

```
[44]: LinearRegression()
```

```
[45]: lm1.intercept_
```

```
[45]: -14382.161315163685
```

```
[46]: lm1.coef_
```

```
[46]: array([ 36.76149419,  3.50153554,  85.32658561, 498.91963877])
```


Hence,

Price = -14382.161315163685 + 36.76149419 x horsepower + 3.50153554 x curb-weight + 85.32658561 x engine-size + 498.91963877 x highway-L/100km

15 Model - Visualization (Simple Linear Regression Model)


```
[47]: # Fitting the simple linear regression  
# Let's visualize the relationship between 'highway-L/100km' and 'price'  
sns.regplot(x = 'highway-L/100km', y = 'price', data = df)
```

```
[47]: <Axes: xlabel='highway-L/100km', ylabel='price'>
```



```
[48]: # This time, visualize the relationship between 'peak-rpm' and 'price'  
sns.regplot(x = 'peak-rpm', y = 'price', data = df)
```

```
[48]: <Axes: xlabel='peak-rpm', ylabel='price'>
```




16 Model - Residual Plot

The residual plot tells us whether the constant variance assumption is violated or not.

If the 'assumptions' are met, then the linear model is appropriate.

```
[49]: # Let's explore the residual plot of 'peak-rpm' and 'price'  
sns.residplot(x= df['peak-rpm'], y = df['price'])
```

```
[49]: <Axes: xlabel='peak-rpm', ylabel='price'>
```



The residual plot above shows that the residuals are not randomly distributed hence it might not be an appropriate linear model.

17 Model - Visualization (Multiple Linear Regression Model)

How do we visualize multiple linear regression model?

One way is by looking at the fit of the model on distribution plot

```
[50]: # Let's predict the Y_hat from our multiple linear regression model lm1
Y_hat = lm1.predict(Z)

# Create the distribution plot,
# The first argument is the variable of the distribution
# The second argument, setting histogram = False
# The third argument, setting the color
# Finally, overlapped the two plots

plot1 = sns.distplot(df['price'], hist= False, color= 'r', label = 'Actual_
↳Value')
plot2 = sns.distplot(Y_hat, hist= False, color= 'b', ax = plot1, label =_
↳'Fitted Value')
```

```
/tmp/ipykernel_450/2046564249.py:10: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``kdeplot`` (an axes-level function for kernel density plots).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
plot1 = sns.distplot(df['price'], hist= False, color= 'r', label = 'Actual Value')
```


```
/tmp/ipykernel_450/2046564249.py:11: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```


Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``kdeplot`` (an axes-level function for kernel density plots).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
plot2 = sns.distplot(Y_hat, hist= False, color= 'b', ax = plot1, label = 'Fitted Value')
```



```
[51]: # Alternatively, we can use the code below since sns update will deprecate ↵  
      ↪ distplot()  
plot1 = sns.kdeplot(df['price'], label = 'Actual Value', color = 'r')  
plot2 = sns.kdeplot(Y_hat, label = 'Fitted Value', ax = plot1, color = 'b')
```

18 Model - Polynomial Regression Model

Sometimes when fitting multiple regression model, the relationship between the independent variables and the dependent variable does not appear linear.

We can fix achieve better fit by using polynomial regression model!

Quadratic - 2nd Order

$$\hat{Y} = a + b_1X + b_2X^2$$

Cubic - 3rd Order

$$\hat{Y} = a + b_1X + b_2X^2 + b_3X^3$$

Higher-Order:

$$Y = a + b_1X + b_2X^2 + b_3X^3 \dots$$

```
[53]: # Let's visualize that by defining a new function 'ploy'

# 1. defines a function named ploy that takes in four arguments
def ploy(model, independent_variable, dependent_variable, Name):
    # Creating 100 evenly spaced values between 15 and 55
    x_new = np.linspace(15, 55, 100)
    # Calculates the corresponding y-values by applying the provided model
    y_new = model(x_new)

    # Creating the plot,
    # '.', dots for actual data points from indepdent_variable and dependent_
    ↪variable
    # '-', line connecting the new x-value and their corresponding y-value
    plt.plot(independent_variable, dependent_variable, '.', x_new, y_new, '-')

    # Adding title and changing the background color
    plt.title('Polynomial Fit with Matplotlib for Price ~ Length')
    ax = plt.gca()
    ax.set_facecolor((0.898, 0.898, 0.898))

    # Customizing X and Y labels
    fig = plt.gcf()
    plt.xlabel(Name)
    plt.ylabel('Price of Cars')

    # Displaying and closing the plot
    plt.show()
    plt.close()

# We can use this function to visualize the distribution plot
```

```
[63]: # Get the variables and then use a polynomial of the 3rd order
x = df['peak-rpm']
y = df['price']
# Fit the polynomial using the function polyfit and poly1d to display the
    ↪function
f = np.polyfit(x,y,3)
p = np.poly1d(f)
print(f)
print(p)
```

```
[-3.66921457e-06  6.03980389e-02 -3.29227169e+02  6.07145229e+05]
      3          2
-3.669e-06 x + 0.0604 x - 329.2 x + 6.071e+05
```

Reference: <https://www.coursera.org/learn/data-analysis-with-python>